

# An SPL compiler in Rust

Ties Klappe  
`s1030293`

April 19, 2022

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Language Choice . . . . .	3
1.2	SPL . . . . .	3
1.2.1	Examples . . . . .	4
<b>2</b>	<b>Lexing &amp; Parsing</b>	<b>5</b>
2.1	Scanner . . . . .	5
2.2	Grammar . . . . .	6
2.2.1	Modifications . . . . .	6
2.2.2	Dangling Else . . . . .	7
2.2.3	Lexer Hack . . . . .	8
2.3	Abstract Syntax Tree . . . . .	8
2.3.1	Expressions . . . . .	8
2.4	Parser . . . . .	9
2.4.1	Example Parser Function . . . . .	10
2.4.2	Chain Rules . . . . .	11
2.5	Pretty Printer . . . . .	11
<b>3</b>	<b>Analyses &amp; Typing</b>	<b>12</b>
3.1	Binding time analysis . . . . .	12
3.1.1	Declaration order . . . . .	12
3.1.2	Global variable restriction . . . . .	13
3.1.3	Context and scopes . . . . .	13
3.2	Return path analysis . . . . .	14
3.3	Type checking . . . . .	15
3.3.1	AST modifications and decorations . . . . .	15
3.3.2	Hindley-Milner: Polymorphic type inference . . . . .	16
	Variable restriction . . . . .	16
	Void restriction . . . . .	17
	Typing variables with fields . . . . .	17
3.3.3	Mutual recursion . . . . .	17
<b>4</b>	<b>Code Generation</b>	<b>18</b>
<b>5</b>	<b>Extension</b>	<b>19</b>
<b>6</b>	<b>Conclusion</b>	<b>20</b>



# Chapter 1

## Introduction

In this report we introduce the compiler for a programming language called SPL, implemented in Rust. This first chapter gives the motivation for choosing Rust as the implementation language, as well as some general information about SPL itself. Chapter 2 covers the frontend of the compiler, i.e. lexing, parsing and pretty printing. Chapter 3 is concerned with semantic analyses and type checking. In Chapter 4, we discuss code generation. In Chapter 5 we propose and describe an extension to the project. Finally, Chapter 6 provides a conclusion and reflection on the project as a whole.

### 1.1 Language Choice

The programming language that was chosen to create the compiler is Rust. Although the teachers advised to use the pure functional programming language Haskell, I was not comfortable enough writing Haskell code for a relatively complex application such as a compiler. Rust was chosen because I have used Rust prior to this course. Besides, it provides some nice (functional) features such as pattern matching and closures, which allows for writing expressive code. It also produces very fast programs, similar to C/C++, while suffering less from issues common to the latter two, such as memory unsafety.

### 1.2 SPL

Simple Programming Language (SPL) is a programming language with a simple C-like syntax. It is a strict language, meaning function arguments must be fully evaluated before a function is called. It is also first-order, as functions may not take other functions as arguments. SPL has a polymorphic type system, where a given function may be called with different parameter types. There are several overloaded operators that work on some or all types, and a number of built-in polymorphic functions. For an overview of the available operators, see Section 2.2.1 and Appendix A.

At the top level, SPL programs consist of a number of variable and function declarations. Users can specify the type of variables explicitly, or declare them as `var`, in which case the type will be inferred. Supported data types are `Int`, `Bool`, `Char`, `[t]` (arrays of type `t`) and `(t1, t2)` (tuples containing values of type `t1` and `t2`).

For function declarations, the function type signature can also be specified or left out. Function signatures consist of zero or more parameter types, as well as a return type. The return

type can also be `Void`, indicating the lack of a return value. Function bodies must contain one or more statements. Available statements are standard `if` (with an optional `else` branch), `while` loops, variable declarations, variable assignments, function calls and returns.

### 1.2.1 Examples

An example in SPL that checks whether a character is a vowel is given in Listing 1.1. An example in SPL that computes the sum of the first  $n$  natural numbers is given in Listing 1.2.

```
1 isVowel(c) :: Char -> Bool
2 {
3   if ( c == 'a' || c == 'e' || c == 'i' || c == 'o' || c == 'u' ) {
4     return True;
5   }
6
7   return False;
8 }
9
10 main() :: -> Void
11 {
12   // True
13   print(isVowel('a'));
14
15   // False
16   print(isVowel('b'));
17 }
```

Listing 1.1: Vowels example

```
1 // Compute the sum of the first n natural numbers.
2 sum(n) :: Int -> Int
3 {
4   var total = 0;
5
6   while(n > 0) {
7     total = total + n;
8
9     n = n - 1;
10  }
11
12  return total;
13 }
14
15 main() :: -> Void
16 {
17   // 15
18   print(sum(5));
19
20   // 5050
21   print(sum(100));
22 }
```

Listing 1.2: Natural numbers example

## Chapter 2

# Lexing & Parsing

The compiler frontend consists of a scanner and a parser. We have chosen the separate scanner/parser approach because it simplifies the parser: instead of having to parse the raw input text, input is first converted into a sequence of abstract tokens. Moreover, this allows us to perform several processing steps such as filtering whitespace, skipping comments and performing character and string escapes before parsing. That way the parser can focus on determining the structure of the input, without having to worry about irrelevant issues like layout.

A side effect of the separate lexer and parser is that both are easy to unit test, and as a result we have a sizeable set of unit tests for both components, which are run using continuous integration in the Git repository.

In this chapter, we describe the scanner, modifications to the provided grammar, the parser and the challenges encountered while developing the frontend of the compiler.

### 2.1 Scanner

The scanner (sometimes also called lexer or tokenizer) splits up an input text into a list of abstract tokens which represent one or more characters of concrete input. The scanner filters the input by removing whitespace characters, single-line comments and multi-line comments. If a multi-line comment was not closed, an appropriate Error token is passed on to the parser. For each token type, there exists an associated regular expression. For example, an identifier is described by the regex `[[:alpha:]](_|[:alnum:]]*)`. This approach allows for a very generic and straightforward implementation of the scanner: most tokens are generated by generic regex matching, and only for literal tokens we have to do some 'manual' work by extracting the literal value and placing it into a literal token. For `Char` and `String` literals this includes removing possible escape sequences. The following characters have to be escaped within both literals: `\`, `\n`, `\r`, `\t`. For a `Char` literal, a single quote must be escaped. For a `String` literal, a double quote must be escaped.

The scanner implements the `Iterator`<sup>1</sup> trait. This allows the parser to *collect* the tokens into a vector of tokens, and ultimately into a `Tokens` object. The `Tokens` type implements all traits that are required by the parser combinator library, which will be introduced in the paragraph 2.4.

---

<sup>1</sup><https://doc.rust-lang.org/std/iter/trait.Iterator.html>

## 2.2 Grammar

Before I continue with the parser, it makes sense to look at the grammar and the modifications to it. The grammar on which the parser is based can be found in Appendix A. There were several issues with the provided grammar that needed to be fixed; these issues were primarily related to left recursion and operator precedence. We use a top-down parsing approach, which means the parser cannot deal with rules like the following (*left recursion*):

```
Field = [Field ( '.' 'hd' | '.' 'tl' | '.' 'fst' | '.' 'snd' )]
```

When parsing sentences of this form, the parser would keep recursively deriving the right-hand `Field`, never finishing. For example, given a token `.hd`, the parser would not be able to determine, in general, whether more `Fields` follow, and so it has to decide to parse `Field` again.

Moreover, consider the following (simplified) expression rule:

```
Exp = Exp '+' Exp
      | Exp '*' Exp
      | num
```

This rule, besides being left-recursive and ambiguous, also does not consider operator precedence. A sentence like `1 + 2 * 3` could be parsed as `(1 + 2) * 3` or `1 + (2 * 3)`, with the former clearly being wrong assuming the usual order of operations. To solve the aforementioned issues, we have modified the provided grammar in several places, which are described in the next paragraph.

### 2.2.1 Modifications

- `FArgs = [ FArgs ',' ] id` has been rewritten to `FArgs = id [ ',' FArgs ]` in order to remove left-recursion.
- `Field = [ Field ( '.' 'hd' | '.' 'tl' | '.' 'fst' | '.' 'snd' ) ]` has been rewritten to  
`Field = ( '.' 'hd' | '.' 'tl' | '.' 'fst' | '.' 'snd' ) [ Field ]`  
 We have done this in order to remove left-recursion and to make field access consistent in the grammar, i.e. optional at the location where it is used instead of optional in its definition.
- Following the previous point, we have made `Field` optional in assignment statements:  
`id [ Field ] '=' Exp ';' ;`  
 Field access is also optional in expression atoms, as discussed in the next point.
- The `Exp` nonterminal has been rewritten to remove left-recursion and to encode operator precedence in the grammar. Each precedence level has a separate rule in the grammar. The precedence levels are listed below, lowest to highest. These were partly inspired by Haskell's operator precedence.

1. `||`
2. `&&`
3. `==, !=, <, <=, >, >=`
4. `:`
5. `+, -` (binary)

6. `*`, `/`, `%`
7. `-` (unary), `!`

Operator fixity is also encoded in the grammar: infix operators always occur between nonterminals, while prefix operators always precede nonterminals. Operators of the same precedence level can be chained. However, associativity is not taken care of in the grammar. Rather, chained operators are handled during parsing, manually enforcing the desired associativity.

At the highest precedence level are expression atoms:

- variables with optional field access;
  - integer, character, boolean and string literals;
  - bracketed expressions;
  - function calls;
  - the empty list literal;
  - tuples.
- `int = [ '-' ] digit+` rewritten to `int = digit+`  
Negative numbers are consistently handled by the unary minus operator, obviating the need for negative-number literals.

### 2.2.2 Dangling Else

A common problem with grammars and parsers of programming languages is the “dangling else” problem, as illustrated in Listing 2.1. Since omitting braces is allowed for if statements where the true/false cases consist of single statements, ambiguity is introduced (does the `else` belong to the outer or inner `if`?). Some languages, like C, solve this by always associating the else branch with the closest if. Other languages, like Rust or Swift, choose to always enforce braces. We have opted for the second solution in the grammar, since it tends to make code more readable, and can even prevent difficult-to-diagnose bugs.<sup>2</sup>

```
1 if (condition1)
2   if (condition2)
3     statement1;
4   else
5     statement2;
```

Listing 2.1: Dangling else

---

<sup>2</sup><https://dwheeler.com/essays/apple-goto-fail.html>



### 2.2.3 Lexer Hack

Some programming languages may have (syntactical) constructs which could have several different meanings. Two examples in C, taken from the lecture slides, are the following:

```
1 void fun()
2 {
3   T (x);
4 }
```

Listing 2.2: Function call or cast?

```
1 void fun()
2 {
3   T * x;
4 }
```

Listing 2.3: Multiplication or pointer declaration?

Here, the scanner cannot know how to tokenise the given program without additional contextual information. Modern compilers usually solve this issue by feeding this contextual information from a later analysis or parsing stage back into the scanner, allowing it to disambiguate the input.

Luckily, SPL does not contain such constructs, which simplifies the scanner and maintains a clean scanner/parser separation.

## 2.3 Abstract Syntax Tree

The abstract syntax tree is a data structure to represent the syntactic structure of programs in an abstract way, without concerning itself with issues like bracketed expressions and layout.

For the definition of the abstract syntax tree, I have taken advantage of Rust’s type system. Rust supports algebraic data types: product types in the form of **structs** and tuples, and sum types in the form of **enums**. Using these constructs, combined with Rust’s **Vec<T>** (growable list) and **Option<T>** (optional, cf. **Maybe** in Haskell), the abstract syntax tree structure is built based on the grammar. For convenience and in order to flatten the structure, the AST slightly deviates from the grammar in several places:

- **FArgs** is directly included in the **FunDecl** type, as it maps one-to-one to a list of identifiers;
- **FTypes** and **RetType** have been “folded” into the **FunType** structure;
- **BasicType** has been folded into the **Type** enumeration;
- The **Exp** operator precedence ladder has been flattened into a single **Expr** type;
- Optionality of variable and function types is represented using **Option<T>**.

### 2.3.1 Expressions

An interesting challenge encountered when defining the expression type is that types in Rust generally must have a size known at compile time. This shows up when defining a recursive data type for expressions. For instance, we would like to represent an expression containing an addition as follows:

```
1 enum Expr {
2   Add(Expr, Expr),
3   ...
4 }
```

However, this results in a recursive type of infinite size, which Rust cannot deal with. The way to solve this is to “box” the contained expressions:

```

1 enum Expr {
2     Add(Box<Expr>, Box<Expr>),
3     ...
4 }

```

In Rust, `Box<T>` is a pointer to heap-allocated data. By doing this, the resulting type *does* have a known size, as the contained expressions are now simply pointers. Although this does incur some (minimal) performance overhead, in practice the difference is negligible. In addition, Rust supports automatic dereferencing in many places, which makes working with these boxed values effectively the same as working with (references to) the values themselves. This also avoids the need for automatic boxing and unboxing.

## 2.4 Parser

For parsing SPL files, we employ a top-down approach. This is achieved using the `nom`<sup>3</sup> parser combinator library for Rust. Parsers created using `nom` are LL( $k$ ), meaning they scan left-to-right and perform leftmost derivation of input sentences, using a lookahead of  $k$  tokens. Parser combinator libraries in general allow building up larger parsers by first writing parsers for specific constructs, and then combining these parsers using combinators. `nom` was chosen because it is the state-of-the-art parser combinator library for Rust, with over 2 million downloads per month. It also strikes a nice balance between usability and efficiency: writing a top-down (e.g. recursive descent) parser by hand would likely result in a faster parser overall, but `nom` is much more user-friendly and flexible in case the grammar changes.

A benefit offered by parser combinators is that the resulting parsers often match their grammar quite closely. Since Rust is more verbose than a language like Haskell, the parser code does not resemble the grammar syntactically, but we have attempted to keep the structure similar to the grammar. Most nonterminals in the grammar have a parser function that can (attempt to) parse the given construct. Higher-level constructs then use combinator functions provided by `nom`, which take one or more parsers and return a new parser function with the desired behavior. For example, there is the `opt` combinator, mapping to the `rule = [ token ]` construct in (e)BNF grammars, which takes a parser `p` and returns a parser that applies `p` zero or one times. A more complete overview of the supported combinators can be found in `nom`'s documentation.<sup>4</sup>

The parser operates on a list of tokens produced by the scanner. It walks the list of tokens from left-to-right while consuming tokens and producing nodes and leaves of the abstract syntax tree data type. The AST is described in more detail in Section 2.3.

Because each parser is a function, I have unit tests for most parser functions. Moreover, we generate a set of integration tests for the parser (and scanner) from the provided set of example SPL programs, which test whether those examples parse correctly and completely.

---

<sup>3</sup><https://lib.rs/crates/nom>

<sup>4</sup>[https://github.com/Geal/nom/blob/main/doc/choosing\\_a\\_combinator.md](https://github.com/Geal/nom/blob/main/doc/choosing_a_combinator.md)

### 2.4.1 Example Parser Function

To give an idea of how parsers in Rust/nom tend to look, Listing 2.4 shows the implementation of a parser for SPL function declarations. The type signature indicates that the function takes a list of tokens and returns nom's `IResult` type. `IResult` is defined as an enum with two variants: `Ok((InputType, OutputType))` and `Err(ErrorType)`. In this case, the input type is `Tokens` and the output type is `FunDecl`, ie. the function returns the rest of the tokens and the parsed function declaration. The function itself consists of the following elements:

1. `map(parser, f)` takes a parser and a function and returns a new parser applying the function to the result of the given parser. Here, the function is defined on lines 17-22, and it constructs a `FunDecl` AST node. `|params| expression` is Rust's syntax for closures/inline functions.
2. `tuple((parser1, parser2, ...))` takes a tuple of parsers and returns a parser that succeeds if it can apply all parsers in order. This parses the following elements:
  - `identifier_parser`, parsing a single identifier;
  - `delimited(first, second, third)`, which applies the first, second and third parser in order and only returns the result of the second: a list of parameters surrounded by parentheses;
  - `separated_list0(sep, parser)` parses zero or more of the second parser, separated by the first: a list of identifiers separated by commas;
  - `fun_decl_type_parser` parses a function type signature;
  - `many1(statement_parser)` parses one or more statements (zero statements produces an error).

```
1 fn fun_decl_parser(tokens: Tokens) -> IResult<Tokens, FunDecl> {
2     map(
3         tuple((
4             identifier_parser,
5             delimited(
6                 opening_paren_parser,
7                 separated_list0(comma_parser, identifier_parser),
8                 closing_paren_parser,
9             ),
10            fun_decl_type_parser,
11            delimited(
12                opening_brace_parser,
13                many1(statement_parser),
14                closing_brace_parser,
15            ),
16        )),
17        |(name, params, fun_type, statements)| FunDecl {
18            name,
19            params,
20            fun_type,
21            statements,
22        },
23    )(tokens)
24 }
```

Listing 2.4: Function declaration parser

### 2.4.2 Chain Rules

As mentioned previously, expressions allow operators of the same precedence level to be chained, e.g.  $1 + 2 + 3$  is legal syntax. For left-associative operators, `nom` supports the `fold_many0` combinator. This operator applies a parser `p` zero or more times, performing a left fold in the process, in a similar fashion to the `pChainl` combinator introduced in the lectures. In this way, we can build up a left-associate tree of operator applications.

Unfortunately, there is no right chain combinator. This means that to parse the list constructor operator `(:)`, the implementation consists of a parser that collects the chain to a vector and then walks the vector right-to-left to build up the abstract syntax tree, instead of a simple right fold.

## 2.5 Pretty Printer

The pretty printer was implemented using the `pretty-trait`<sup>5</sup> library. This library provides functionalities to print arbitrary types by joining constructs, adding indentations and newline symbols and grouping related content together. The `to_pretty` function was implemented for every type of the AST. Pretty printing an entire AST can be done by simply converting the top level AST object (*Program*) to a `Pretty` object and printing it.

When printing expressions, some parts may need to be parenthesized. This is the case when a child expression node in the AST has a lower precedence than its parent: the child expression must be parenthesized in this case. For example, figure 2.1 displays an AST where the addition operator occurs below the multiplication operator. Since addition has lower precedence than multiplication (see section 2.2.1), the addition expression must be parenthesized when we pretty print this AST:  $(1 + 2) * 3$ . Figure 2.2 displays an AST where no child expressions have lower precedence, hence no parenthesis are needed when we pretty print this AST:  $1 + 2 * 3$ .




Figure 2.1: Parentheses required




Figure 2.2: No parentheses required

The pretty printer does not keep comments, newlines and meaningless white spaces. Tabs have a size of four white spaces. A single line has a maximum length of forty white spaces.

Finally, several tests are written to verify that the AST of the original input file and the AST of the pretty printed file are identical.

---

<sup>5</sup><https://lib.rs/crates/pretty-trait>

## Chapter 3

# Analyses & Typing

The middle end of the compiler implements binding time analysis, return path analysis and polymorphic type inference. In this chapter the implementation details of these functionalities are described. Several pieces of Rust code are shown to substantiate the statements made for each analysis, which are copied directly from the source code (available on Gitlab).

### 3.1 Binding time analysis

The binding time analysis is largely integrated in the type inference process, on which paragraph 3.3 elaborates. This paragraph explains all the rules on scoping and identifiers that are enforced by the compiler.

#### 3.1.1 Declaration order

A problem arises when a programmer calls a function  $f$  that has not yet been defined. The AST constructed from the tokens in the previous phase simply reads the source code from top to bottom, and thus the AST nodes are not sorted on usage thereafter.

One of the possible solutions to this problem is to construct the call graph from the AST, performing topological sort and resorting the AST afterwards. In such a call graph, functions are presented as nodes and function calls as edges. As global variables are not allowed to make function calls (see section 3.1.2), we can effectively move all global variables to the front of the declarations list of the AST and sort them among themselves (they may be defined in terms of each other). Finally, we also need to reorder local variables inside function declarations. The Reorder procedure below provides an abstract overview of the complete AST reordering.

Reorder( $T$ )

1. Construct call graph  $G_v$  of global variables of  $T$
2. Run topological sort on  $G_v$
3. Insert global variable declarations from  $G_v$  at the front of  $T'$
5. Construct call graph  $G_f$  of functions of  $T$
6. Run topological sort on  $G_f$
7. Insert function declarations from  $G_f$  at the back of  $T'$
8. Foreach function declaration  $fd$ , apply Reorder( $fd$ ) for local variable resorting
9. Return  $T'$

### 3.1.2 Global variable restriction

A global variable is restricted from using function calls in the expressions representing its value. Allowing this would impose a difficulty when defining semantics for programs such as the one given in listing 3.1. Here, we don't know the value for  $x$ , because its expression is defined by a function  $f$  which is defined in terms of  $x$ . Using global variables in functions is in general not a problem, as can be seen in listing 3.2. Another solution would have been to analyse all function calls made by global variables, determine whether that function is defined in terms of that same global variable and give a compilation error if that is the case.

```
1 var x = f();
2 f() {
3   return x;
4 }
```

Listing 3.1: Global variable restriction

```
1 var x = 1;
2 f() {
3   return x;
4 }
```

Listing 3.2: Global variable restriction

### 3.1.3 Context and scopes

The context that is used throughout type checking, is composed of at least two type environments  $\{functions, [vars]\}$ . A type environment is implemented as a hashmap, that maps an identifier to a type scheme. The environments that compose the context have fairly self-explanatory names: the *functions* environment contains all function identifiers and types, whereas the *[vars]* environment contains a list of type environments for the identifiers of variables at a given moment in the algorithm. This list of environments always contains the environment for global variables at index zero. As the algorithm runs, environments for new scopes (function arguments and local variables) are pushed to the list, and popped when the algorithm leaves that scope.

The compiler enforces the following rules:

1. Function and variable names may overlap. For example, a program that has a function named  $f$  and a global variable also named  $f$  is fine

As SPL only has first order functions, it is nice feature to allow the same name for a function and a variable.

2. Variables may be shadowed, but a warning is produced

If a global variable  $x$  exists, and a local variable or function parameter is also identified by  $x$ , the last defined  $x$  will be used. Listing 3.3 shows a small example: if  $x$  occurs in the remaining body of  $f$ , the  $x$  declared at line 4 is the variable being referred to. Two warnings will be given: the first one when the function parameter at line 3 hides the global variable  $x$ , and the second one when local variable  $x$  at line 4 hides both the function parameter and global variable  $x$ .

3. Variable names must be unique in the same scope

As variables may be declared in a different order than in which they are used, allowing redefinition of variables in the same scope would make it very hard to determine which variable an identifier is referring to. For that reason, all variable names inside the same scope must be unique. A small example is given in listing 3.4: trying to compile this code would fail with an error that function argument  $x$  was previously declared.

```

1 var x = 10;
2
3 f(x) {
4     var x = 1;
5     ...
6 }

```

Listing 3.3: Local variable shadowing

```

1 var x = 10;
2
3 // This is not allowed.
4 f(x, x) {
5     ...
6 }

```

Listing 3.4: Function parameter duplicate

## 3.2 Return path analysis

Return path analysis guarantees that a function which is supposed to return something actually does so. This prevents executing code where for example the true branch of an if-statement returns a value, and the false branch does not return a value. In languages such as C, assigning the result of the false branch of this example to a variable simply results in undefined behavior.

Return path analysis is performed immediately before we run the type checker on a function declaration. In the implementation, a trait named **ReturnPathAnalysis** is used and implemented for the AST nodes **Statement** and **FunDecl**. This trait provides just a single function: **check\_returns**, which returns information about the return type of the function. This information is one of four variants:

1. **None**: the function did not return anything, this results in an error.
2. **Incomplete**: the function did not return something on all branches. This is the case when for example the true branch of an if-statement did return a value, and the false branch did not (like in the example above), and no other returns occur inside the function body.
3. **Some**: the function did return something on all branches. We are not interested in the type, as this will be inferred directly after this analysis by the type checker.
4. **Void**: the function returns on all branches, but no expression is returned which means that the return type is void.

The implementation of **check\_returns** for a **FunDecl** AST node is made by a fairly simple traversal of the statements inside the function body, the most interesting part is given in listing 3.5. The implementation for a vector of statements simply loops over all statements, until a complete return (**Some** or **Void**) is encountered after which it immediately returns. Otherwise the function returns **None** or **Incomplete**, depending on the results of the recursive calls.

```

1 fn check_returns(&self) -> ReturnType {
2     match self {
3         Statement::If(i) => match (i.if_true.check_returns(), i.if_false.check_returns()) {
4             (ReturnType::None, ReturnType::None) => ReturnType::None,
5             (ReturnType::Void, ReturnType::Void) => ReturnType::Void,
6             (ReturnType::Some, ReturnType::Some) => ReturnType::Some,
7
8             // Here we have a type conflict, this will become a proper error from the
9             // type inferencer in the next analysis.
10            (ReturnType::Some, ReturnType::Void) => ReturnType::Some,
11            (ReturnType::Void, ReturnType::Some) => ReturnType::Some,
12        },
13        Statement::While(w) => w.body.check_returns(),
14        Statement::Return(e) => match e {
15            Some(_) => ReturnType::Some,
16            None => ReturnType::Void,
17        },
18        _ => ReturnType::None,
19    }
20 }

```

Listing 3.5: Return path analysis

## 3.3 Type checking

This section elaborates on the implementation of type checking/inference.

### 3.3.1 AST modifications and decorations

The AST that was produced in the previous phase of the compiler contained different data types for function and variable types. In order to simplify the type checking analysis, these types were merged into a single datatype given in listing 3.6. The parser is modified accordingly, to prevent `Void` to be a valid type for a variable or function parameter.

```

1 pub enum Type {
2     Int,
3     Bool,
4     Char,
5     Void,
6     Function(Vec<Type>, Box<Type>),
7     Tuple(Box<Type>, Box<Type>),
8     List(Box<Type>),
9     Var(Id),
10 }

```

Listing 3.6: SPL types

Types for functions and variables make use of Rust’s *Option* type (similar to *Maybe* in Haskell). Whenever a type is inferred, the AST is updated with *Some(inferred type)* (if no type was specified by the programmer in the code), or the unification of the programmer-specified type and the inferred type applied on the specified type.

The AST nodes are not decorated with source information (yet). Although this is a nice-to-have in order to give errors including the location in the source code, I had to work on other parts of the compiler due to time constraints.



### 3.3.2 Hindley-Milner: Polymorphic type inference

The compiler implements polymorphic type inference, closely following both the lecture slides and the implementation by Martin Gradmüller [1] (W-algorithm). The core of the implementation is made up by two traits: `TypeInstance` and `TypeInference`. The `TypeInstance` trait provides the two functions `ftv` and `apply`, as given in listing 3.7. This trait is implemented for `Type`, `TypeScheme` and `TypeEnv`, allowing to determine the free type variables and to apply substitutions to these datatypes. The `TypeInference` trait provides only a single function `infer` and is implemented for most AST nodes, as given in listing 3.8.

```
1 trait TypeInstance {
2     // Determines the free type variables of a type.
3     fn ftv(&self) -> HashSet<TypeVar>;
4
5     // Apply a substitution.
6     fn apply(&self, subst: &Subst) -> Self;
7 }
```

Listing 3.7: TypeInstance trait

```
1 trait TypeInference {
2     fn infer(
3         &mut self,
4         context: &Context,
5         expected: &Type,
6         generator: &mut VarGenerator,
7     ) -> Result<Subst, String>;
8 }
```

Listing 3.8: TypeInference trait

The `infer` function implements the presented function  $M$  from the lecture slides. The implementations make a clear distinction between inferring the type of statements (for which we normally don't use the expected type) and expressions (for which we unify the inferred and expected type).

Running the complete type inference process is done by creating an empty context in which the builtin functions `print`, `isEmpty` and field operators are loaded, reordering the AST and then calling `infer` on the (mutable) root node of the AST. We then get an empty result (for success) in which the passed AST is completely typed, or an error.

#### Variable restriction

As presenting in the lectures, a program such as the one in listing 3.9 would successfully pass the type checker if we blindly follow the generalization rule for let expressions. As SPL has only the empty list as a polymorph value, we can use *variable restriction* (not generalizing variables) so that the type checker does not accept this problematic program.

```
1 problematic() {
2     var x = [];    // Context: {Aa. x: [a]}, Values: {x = []}
3     x = 3 : x;    // Context: {Aa. x: [a]}, Values: {x = [3]}
4     x = True : x; // Context: {Aa. x: [a]}, Values: {x = [True, 3]}
5 }
```

Listing 3.9: Polymorph values

## Void restriction

Void can be seen as a special type. It is only allowed to be used in return types of functions. This introduces some problems, for example, an expression such as `print(1) == print(2)` effectively compares two voids and should thus not be allowed. Similarly, assigning void to a variable by for example the expression `var x = print(1);` should not be allowed. These examples have one commonality, namely that they use *function calls in expressions* to create a void expression. To disallow such problematic expressions, using function calls that return void are disallowed in the implemented compiler. As this rule does not apply for function calls in statements, writing `print(1);` in a function body is still perfectly fine for the compiler.

## Typing variables with fields

All field operators (`hd`, `tl`, `fst` and `snd`) are treated as builtin unary functions. This means that whenever we want to type an expression such as `x.hd.fst.tl`, we simply convert this expression to an equivalent sequence of function calls whose type we want to infer. In the case of the above example, the we return the type of the expression `tl(fst(hd(x)))`. Applying standard unification and substitution gives the type  $[[a], b] \rightarrow [a]$  of the composed functions, and thus after applying it on `x` the expression has type `[a]` (`a` possibly substituted for its type in `x`).

### 3.3.3 Mutual recursion

As explained in section 3.1.1, the declaration order of functions and variables does not have to be the same as the order in which they are used. A new problem arises when functions are defined in terms of each other: mutual recursion. A small example is given in listing 3.10: function `foo` calls function `bar` and vice versa. When we apply the type inference algorithm as normal, the function identifier `bar` will be unknown when we try to type `foo`, and an error is returned.

```
1 foo() {  
2   return bar();  
3 }  
4 bar() {  
5   var res = foo();  
6   return res + 1;  
7 }
```

Listing 3.10: Mutual recursion example

A solution to this problem, as was presented during the lectures on typing, is to type these functions together. Moreover, we can determine which groups of functions have to be typed together by performing strongly connected component analysis on the call graph. The presented Tarjan algorithm[2] does exactly that, and also performs topological sorting as a byproduct which solves the problems described in section 3.1.1.

Currently, the compiler does not support mutual recursion yet. The compiler does however construct the call graph from the AST by letting the functions represent the nodes and the function calls edges. The `Petgraph`<sup>1</sup> crate is used to construct this graph, and also provides an implementation of Tarjans algorithm. So as a final step for mutual recursion, the compiler does not only need to resort the declarations in the AST, but also type the components together.

---

<sup>1</sup><https://docs.rs/crate/petgraph/latest>

## Chapter 4

# Code Generation

## Chapter 5

## Extension

## Chapter 6

## Conclusion

# Bibliography

- [1] GRABMÜLLER, M. Algorithm w step by step, 2006.
- [2] TARJAN, R. Depth-first search and linear graph algorithms. *SIAM journal on computing* 1, 2 (1972), 146–160.

# Appendix A

## Grammar

```
SPL      = Decl+
Decl     = VarDecl
        | FunDecl
VarDecl  = ('var' | Type) id '=' Exp ';'
FunDecl  = id '(' [ FArgs ] ')' [ '::' FunType ] '{' VarDecl* Stmt+ '}'
RetType  = Type
        | 'Void'
FunType  = [ FTypes ] '->' RetType
FTypes   = Type [ FTypes ]
Type     = BasicType
        | '(' Type ',' Type ')'
        | '[' Type ']'
        | id
BasicType = 'Int'
        | 'Bool'
        | 'Char'
FArgs     = id [ ',' FArgs ]
Stmt      = 'if' '(' Exp ')' '{' Stmt* '}' [ 'else' '{' Stmt* '}' ]
        | 'while' '(' Exp ')' '{' Stmt* '}'
        | id [ Field ] '=' Exp ';'
        | FunCall ';'
        | 'return' [ Exp ] ';'
Exp       = Disjunc
Disjunc   = Conjunc ( '||' Conjunc )*
Conjunc   = Compare ( '&&' Compare )*
Compare   = Concat ( ( '==' | '!=' | '<' | '<=' | '>' | '>=' ) Concat )*
Concat    = Term ( ':' Term )*
Term      = Factor ( ( '+' | '-' ) Factor )*
Factor    = Unary ( ( '*' | '/' | '%' ) Unary )*
Unary     = ( '-' | '!' ) Unary
        | Atom
Atom      = id [ Field ]
        | int | char | string | 'True' | 'False'
        | '(' Exp ')'
```

```

        | FunCall
        | '['
        | '(' Exp ',' Exp ')'
Field    = ( '.' 'hd' | '.' 'tl' | '.' 'fst' | '.' 'snd' ) [ Field ]
FunCall  = id '(' [ ActArgs ] ')'
ActArgs  = Exp [ ',' ActArgs ]
int      = digit+
id       = alpha ( '_' | alphaNum)*

```

- **char** is a `'`, followed by a char escape sequence or any Unicode character besides `'` and `\`, followed by `'`.
- a char escape sequence is a `\` followed by `'`, `\`, `n`, `r` or `t`.
- **string** is a `"`, followed by zero or more string escape sequences or any Unicode character besides `"` and `\`, followed by `"`.
- a string escape sequence is a `\` followed by `"`, `\`, `n`, `r` or `t`.